

Assignment Project Exam Help

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro

SUPPORT VECTOR MACHINE

Mainly based on

<https://nlp.stanford.edu/IR-book/pdf/15svm.pdf>

Overview

- SVM is a huge topic
 - Integration of MMDS, IIR, and Andrew Moore's slides here
- Our foci: **Assignment Project Exam Help**
 - Geometric interpretation in <https://eduassistpro.github.io/>
 - Alternative interpretation from a geometric perspective. Add WeChat **edu_assist_pro**
 - Understand the final solution (in the dual form)
 - Generalizes well to Kernel functions
 - SVM + Kernels

$$\mathbf{w}^T \mathbf{x}_i + b = 0$$

Linear classifiers: Which Hyperplane?



- Lots of possible solutions for a, b, c .
- Some methods find a separating hyperplane, but **Assignment Project Exam Help** [according to some criterion goodness]

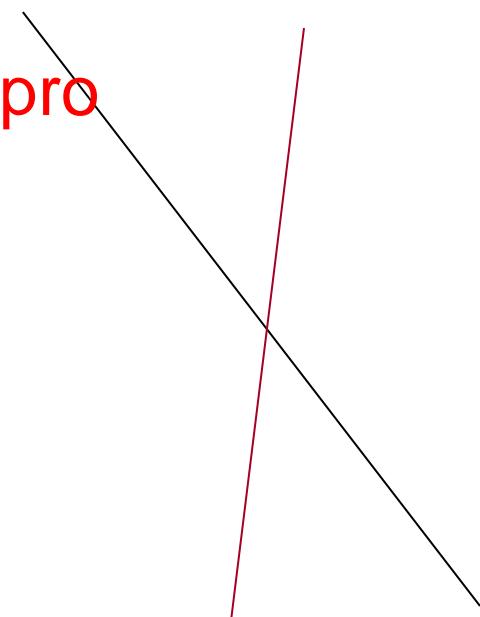
This line represents the decision boundary:

$$ax + by - c = 0$$

- E.g., perceptron
- Support Vector Machine (SVM) finds optimal* solution.
 - Maximizes the distance between the hyperplane and the “difficult points” close to decision boundary
 - One intuition: if there are no points near the decision surface, then there are no very uncertain classification decisions

<https://eduassistpro.github.io/>

Add WeChat **edu_assist_pro**



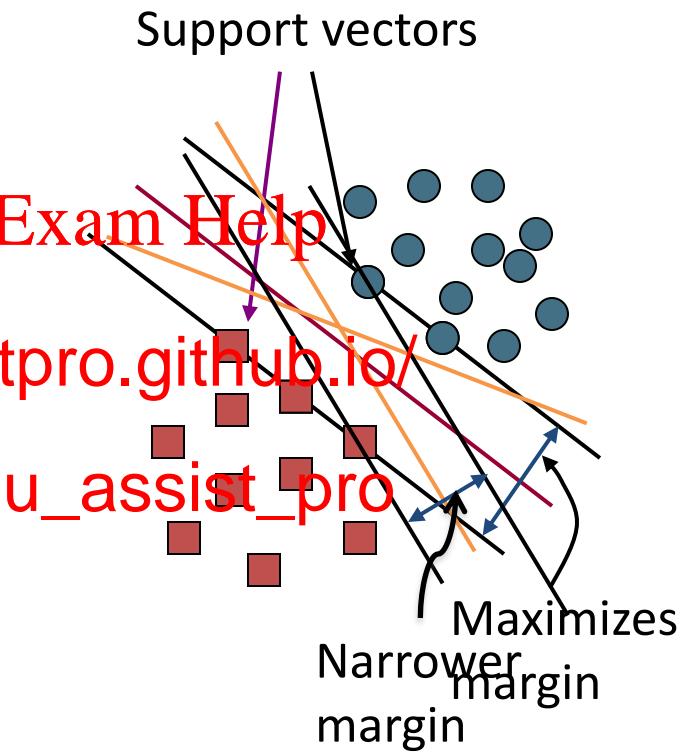
Another intuition

- If you have to place a fat separator between classes, you have less choices, and so the capacity of the model has been decreased



Support Vector Machine (SVM)

- SVMs maximize the *margin* around the separating hyperplane.
 - A.k.a. Large margin classifiers
- The decision function is specified by a subset of samples, the support vectors.
- Solving SVMs is a *quadratic programming* problem
- Seen by many as the most successful current text classification method*



*but other discriminative methods often perform very similarly

Maximum Margin: Formalization

- \mathbf{w} : decision hyperplane normal vector
- \mathbf{x}_i : data point i
- y_i : class of d
- Classifier is: <https://eduassistpro.github.io/>
- Functional margin of w_i is:
Add WeChat edu_assist_pro
 - But note that we can increase this margin simply by scaling \mathbf{w} , \mathbf{b}
- Functional margin of dataset is twice the minimum functional margin for any point
 - The factor of 2 comes from measuring the whole width of the margin

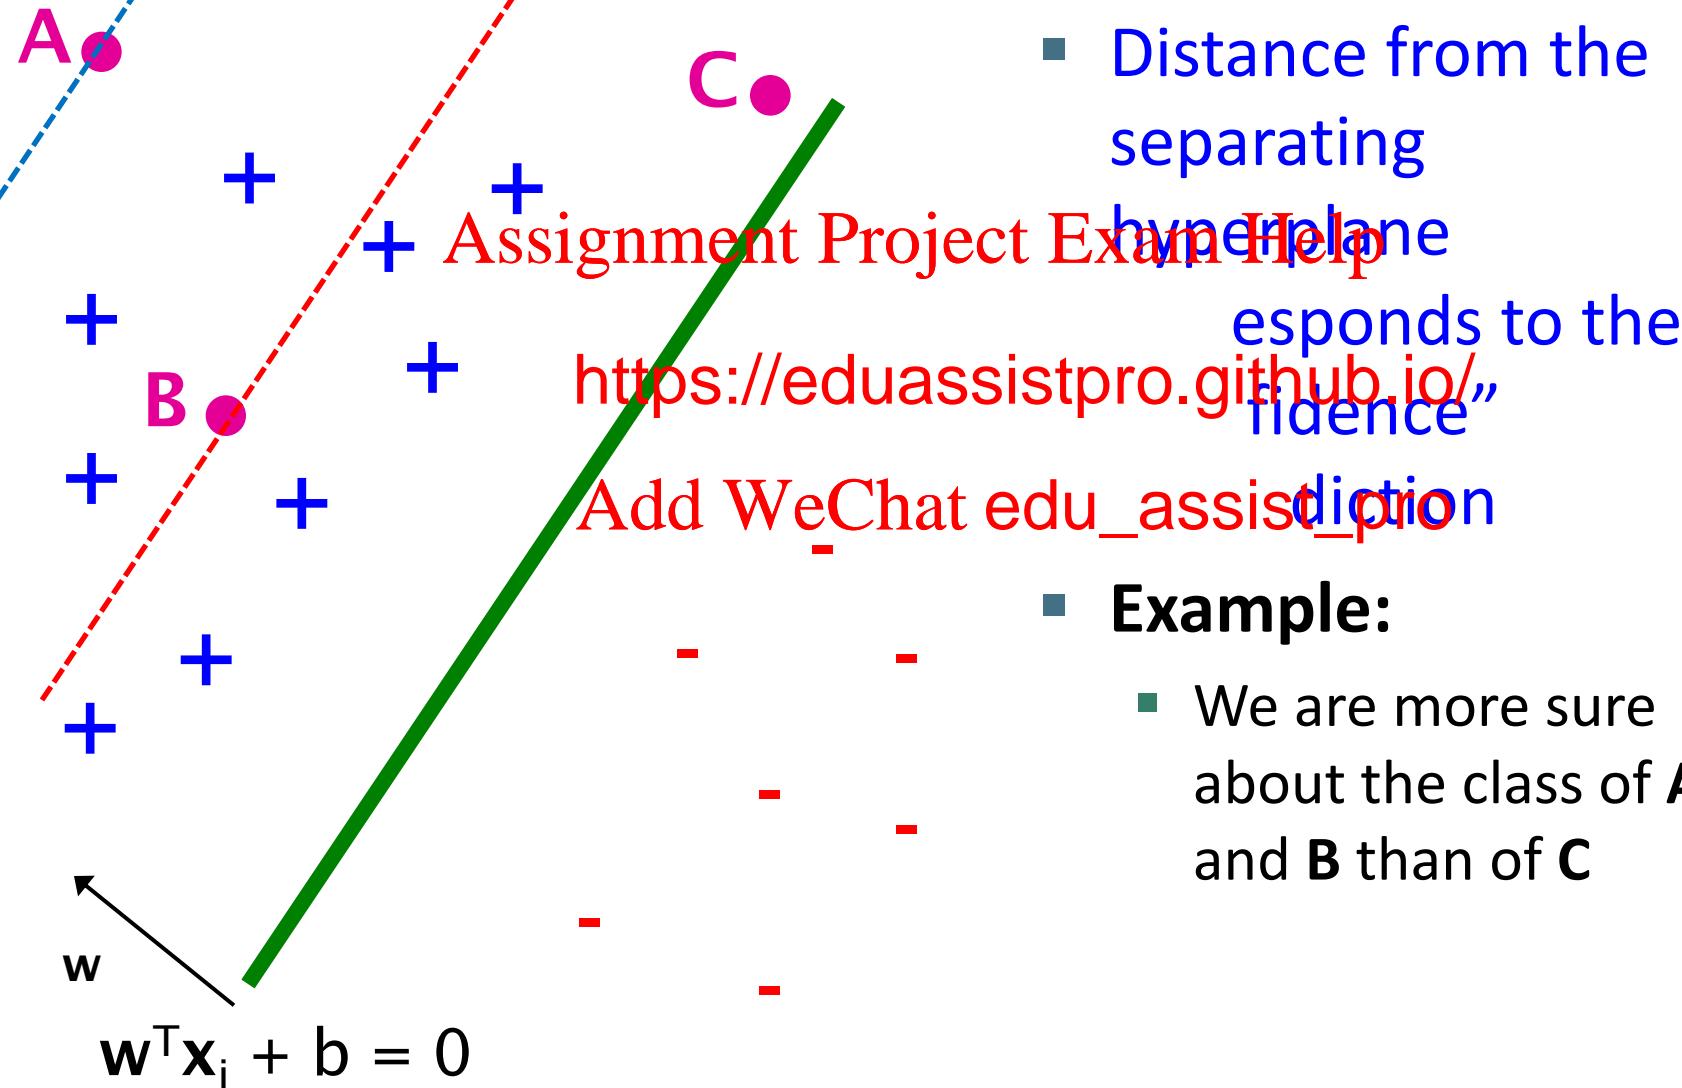
NB: Not 1/0

NB: a common trick

$$\mathbf{w}^T \mathbf{x}_i + b = 7.4$$

$$\mathbf{w}^T \mathbf{x}_i + b = 3.7$$

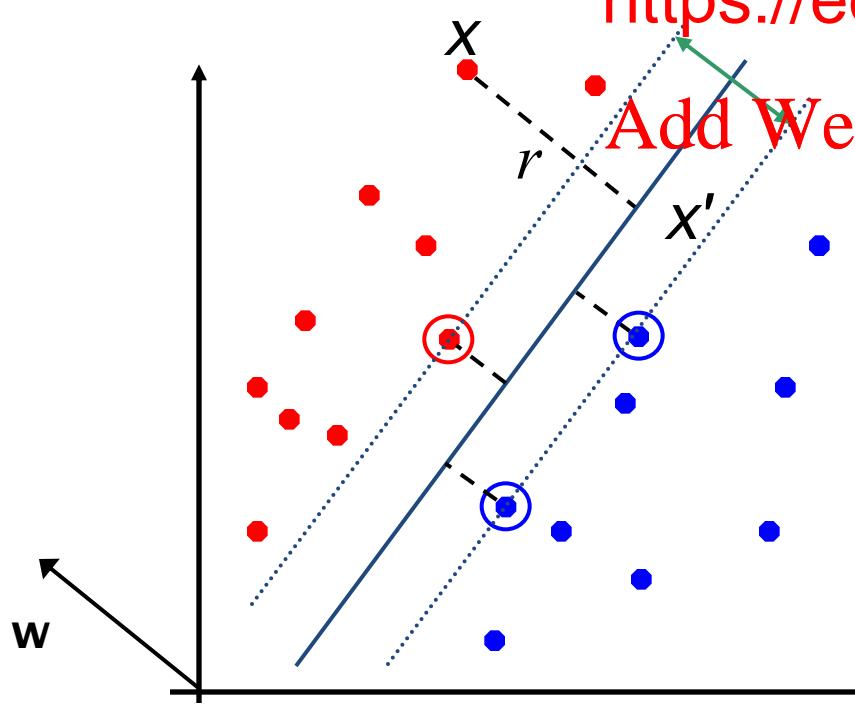
Largest Margin



Geometric Margin

- Distance from example to the separator is
- Examples closest to the hyperplane are *support vectors*.
- Margin ρ* of the separator between support vectors of classes.

$$r = y \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$$



Assignment Project Exam Help

<https://eduassistpro.github.io/>

ic der

ding r :

x' is perpendicular to boundary so parallel to w .

Unit vector is $w/\|w\|$, so line is $rw/\|w\|$, for some r .

$x' = x - yrw/\|w\|$.

x' satisfies $w^T x' + b = 0$.

So $w^T(x - yrw/\|w\|) + b = 0$

Recall that $|w| = \sqrt{w^T w}$.

So $w^T x - yr|w| + b = 0$

So, solving for r gives:

$r = y(w^T x + b)/|w|$

Help from Inner Product

- Remember: Dot product / inner product

Assignment Project Exam Help

$$\|A\| \|B\| \cos \theta$$

<https://eduassistpro.github.io/>

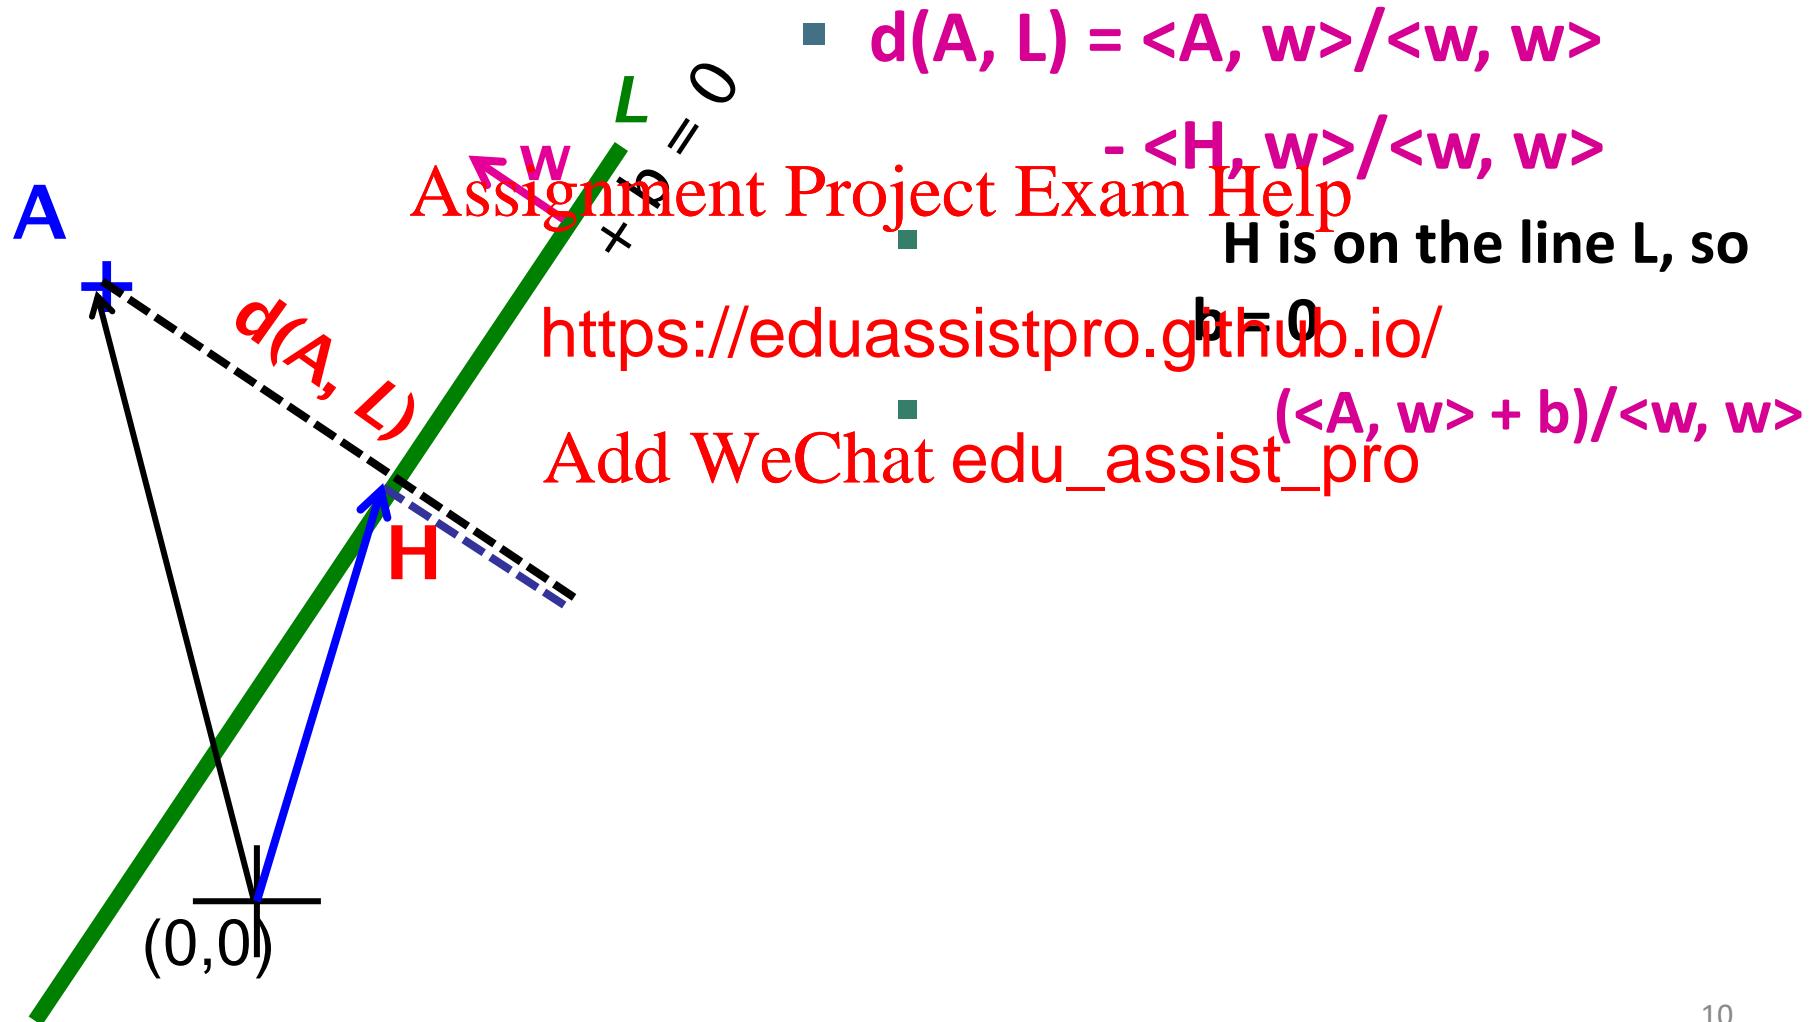
Add WeChat edu_assist_pro

$$\|A\| \cos \theta$$

- (Or use an algebraic derivation similar to prev slide)

- A's projection onto B = $\langle A, B \rangle / \langle B, B \rangle$

Derivation of the Geometric Margin



Linear SVM Mathematically

The linearly separable case

- Assume that all data is at least distance 1 from the hyperplane, then the following two constraints follow for a training set $\{(\mathbf{x}_i, y_i)\}$

Assignment Project Exam Help

$$\mathbf{w}^T \mathbf{x} + b \geq 1 \quad \text{if } y = 1$$

<https://eduassistpro.github.io/>

- For support vectors, the inequality becomes
- Then, since each example's distance from the plane is

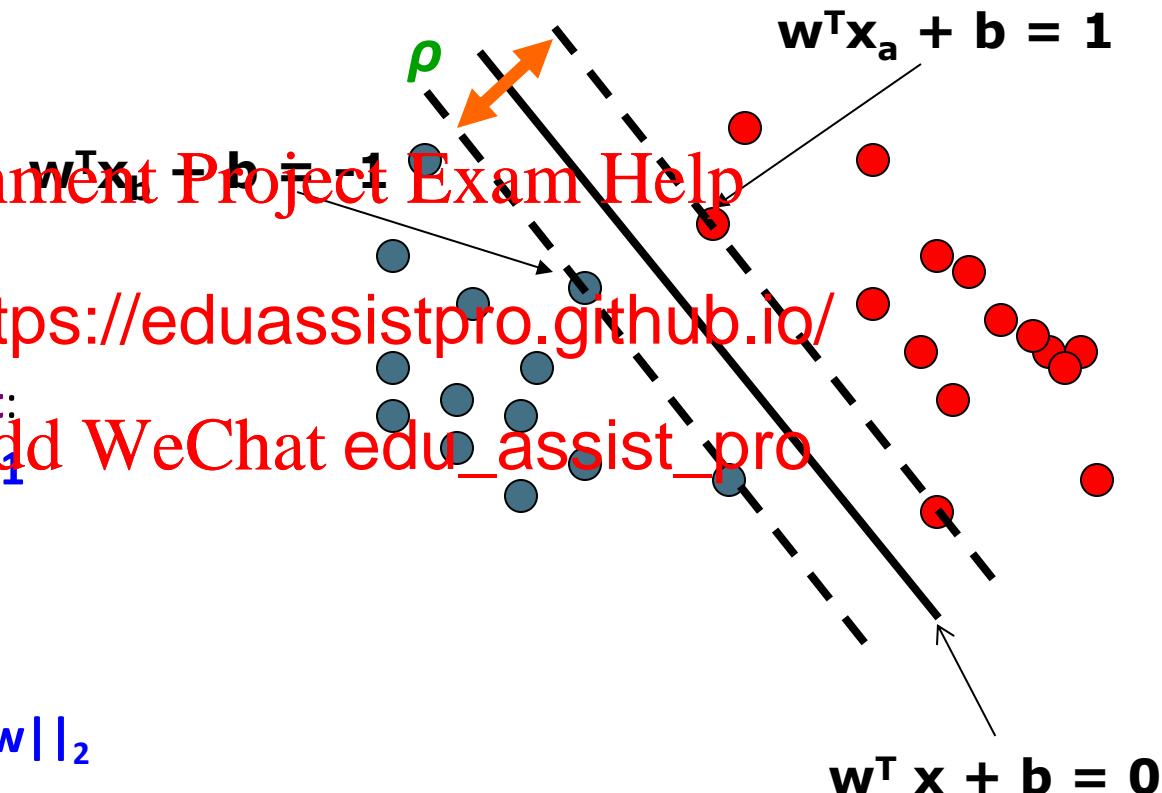
$$r = y \frac{\mathbf{w}^T \mathbf{x} + b}{\|\mathbf{w}\|}$$

- The margin is:

$$\rho = \frac{2}{\|\mathbf{w}\|}$$

Derivation of ρ

- Hyperplane $w^T x + b = 0$
- Extra scale constraint:
 $\min_{i=1,\dots,n} |w^T x_i + b| = 1$
- This implies:
 $w^T(x_a - x_b) = 2$
 $\rho = ||x_a - x_b||_2 = 2/\|w\|_2$



Solving the Optimization Problem

Find \mathbf{w} and b such that

$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$ is minimized;

and for all $\{(x_i, y_i)\}$: $y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$

Assignment Project Exam Help

- This is now optimization problem subject to *linear constraints*
- Quadratic optimization problem, and many (internal) solvers exist for solving them (with many special ones built for SVM)
- The solution involves constructing a *dual problem* where a *Lagrange multiplier* α_i is associated with every constraint in the primary problem:

Find $\alpha_1 \dots \alpha_N$ such that

$Q(\alpha) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ is maximized and

$$(1) \quad \sum \alpha_i y_i = 0$$

$$(2) \quad \alpha_i \geq 0 \text{ for all } \alpha_i$$

Geometric Interpretation

Find \mathbf{w} and b such that

$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$ is minimized;

and for all $\{(x_i, y_i)\}$: $y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$

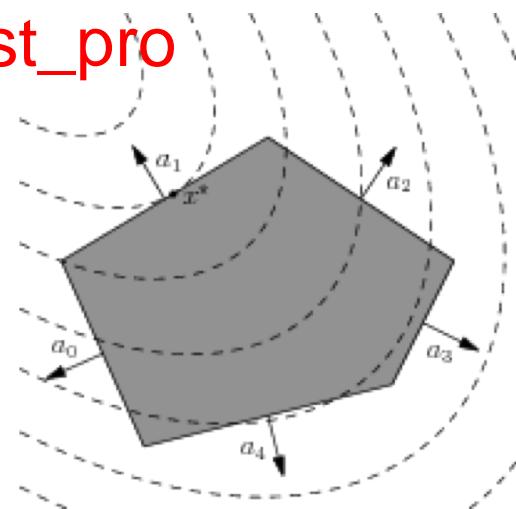
- What are fixe

- Linear constr

- Quadratic objective function:

<https://eduassistpro.github.io/>

Add WeChat edu_assist_pro



The Optimization Problem Solution

- The solution has the form:

$$\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i \quad b = y_k - \mathbf{w}^T \mathbf{x}_k \text{ for any } \mathbf{x}_k \text{ such that } \alpha_k \neq 0$$

- Each non-zero α_i indicates that corresponding \mathbf{x}_i is a support vector.
- Then the scoring function

$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

Assignment Project Exam Help
<https://eduassistpro.github.io/>

$$y_v \cdot \alpha_{sv} \cdot \langle \mathbf{x}_{sv}, \mathbf{x} \rangle + b$$

Add WeChat edu_assist_pro

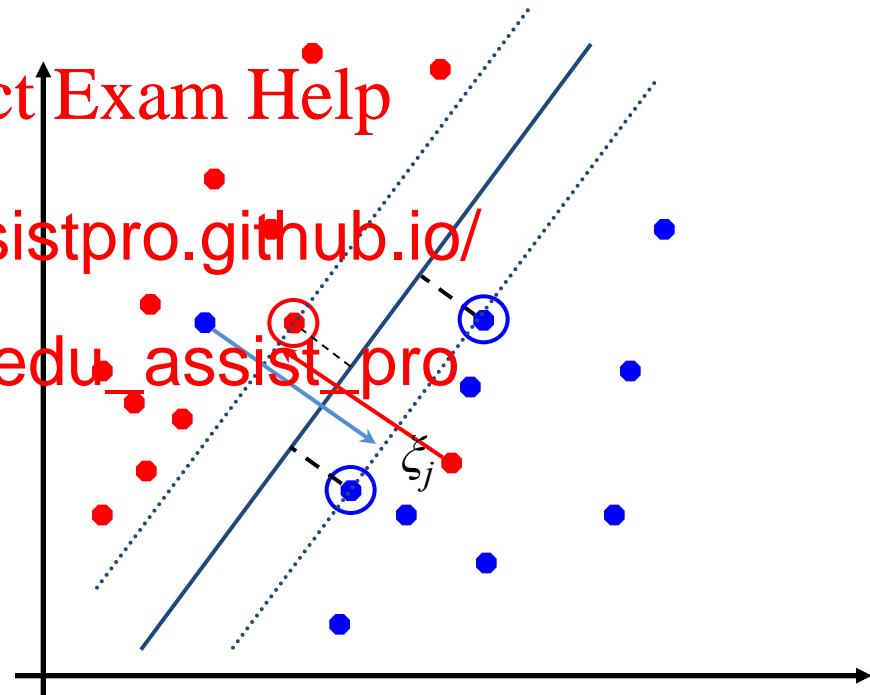
Q: What are the model parameters? What does $f(x)$ mean intuitively?

- Classification is based on the sign of $f(\mathbf{x})$
- Notice that it relies on an **inner product** between the test point \mathbf{x} and the support vectors \mathbf{x}_i
 - We will return to this later.
- Also *keep in mind* that solving the optimization problem involved computing the inner products $\langle \mathbf{x}_i, \mathbf{x}_j \rangle$ between all pairs of training points.

Soft Margin Classification

- If the training data is not linearly separable, *slack variables* & can be added to allow misclass difficult or noi
- Allow some er
 - Let some points be moved to where they belong, at a cost
- Still, try to minimize training set errors, and to place hyperplane “far” from each class (large margin)

Assignment Project Exam Help
https://eduassistpro.github.io/
Add WeChat edu_assist_pro



Soft Margin Classification

[Optional]

Mathematically

- The old formulation:

Find \mathbf{w} and b such that

$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w}$ is minimized and for all $\{(\mathbf{x}_i, y_i)\}$

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1$$

<https://eduassistpro.github.io/>

- The new formulation:

Add WeChat `edu_assist_pro`

Find \mathbf{w} and b such that

$\Phi(\mathbf{w}) = \frac{1}{2} \mathbf{w}^T \mathbf{w} + C \sum \xi_i$ is minimized and for all $\{(\mathbf{x}_i, y_i)\}$

$$y_i (\mathbf{w}^T \mathbf{x}_i + b) \geq 1 - \xi_i \quad \text{and} \quad \xi_i \geq 0 \text{ for all } i$$

- Parameter C can be viewed as a way to control overfitting
 - A regularization term

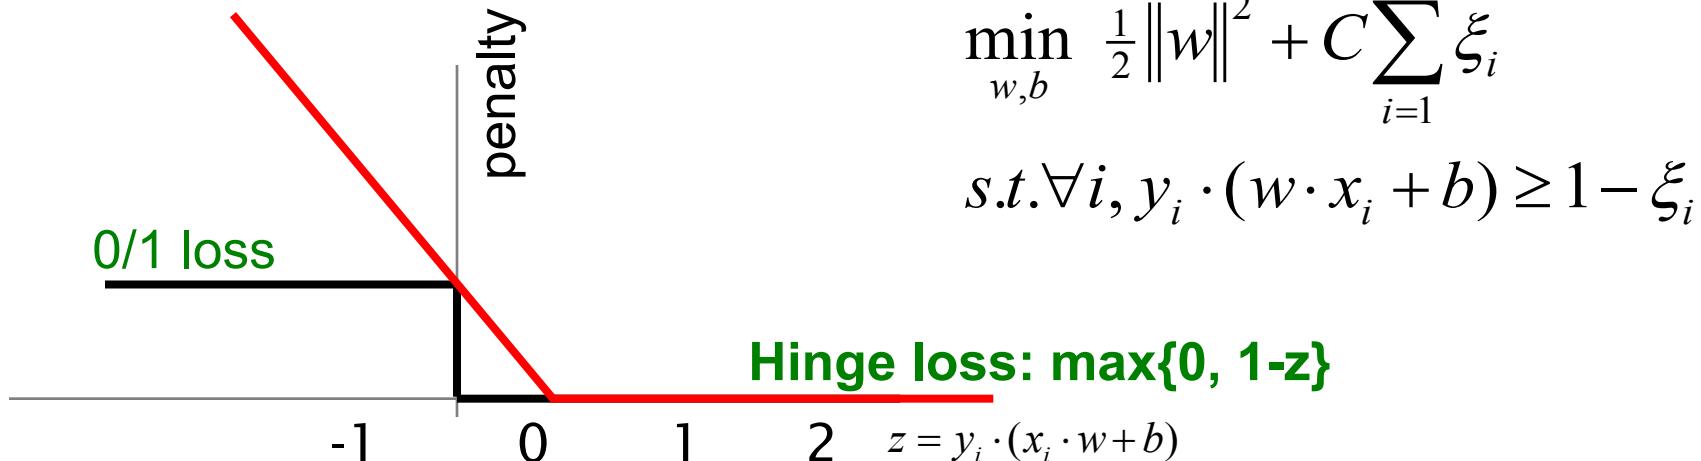
Alternative View

- SVM in the “natural” form

$$\arg \min_{w,b} \frac{1}{2} w \cdot w + C \cdot \sum_{i=1}^n \max\{0, 1 - y_i(w \cdot x_i + b)\}$$

Margin Assignment Project Exam Help L (how well we fit training data)
Hyper-para <https://eduassistpro.github.io/>

- SVM uses “Hinge Loss”:



[Optional]

Soft Margin Classification – Solution

- The dual problem for soft margin classification:

Find $\alpha_1 \dots \alpha_N$ such that

$Q(\alpha) = \sum \alpha_i - \gamma_2 \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ is maximized and

$$(1) \sum \alpha_i y_i = 0$$

$$(2) 0 \leq \alpha_i \leq C \text{ for } \textcolor{red}{\text{https://eduassistpro.github.io/}}$$

- Neither slack variables ξ_i nor their Lagrangian multipliers appear in the dual problem!
- Again, \mathbf{x}_i with non-zero α_i will be support vectors.
- Solution to the dual problem is:

$$\mathbf{w} = \sum \alpha_i y_i \mathbf{x}_i$$

$$b = y_k(1 - \xi_k) - \mathbf{w}^T \mathbf{x}_k \text{ where } k = \operatorname{argmax}_{k'} \alpha_k,$$

w is not needed explicitly for classification!

$$f(\mathbf{x}) = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$$

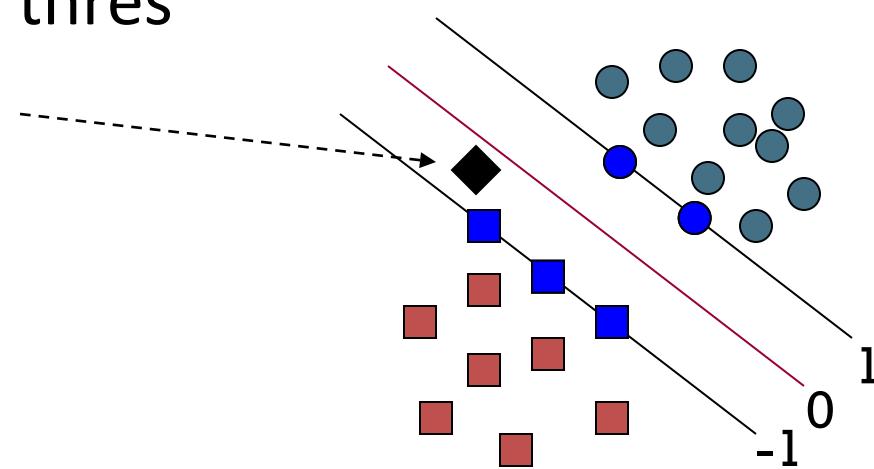
Classification with SVMs

- Given a new point \mathbf{x} , we can score its projection onto the hyperplane normal:
 - i.e., compute score: $\mathbf{w}^T \mathbf{x} + b = \sum \alpha_i y_i \mathbf{x}_i^T \mathbf{x} + b$
 - Decide clas <https://eduassistpro.github.io/>
 - Can set confidence thres

Score $> t$: yes

Score $< -t$: no

Else: don't know



Linear SVMs: Summary

- The classifier is a *separating hyperplane*.
- The most “important” training points are the support vectors; they define the hyperplane
- Quadratic optimization problem: find training points \mathbf{x}_i are support vectors with $\alpha_i > 0$.
- Both in the dual formulation of the problem: the solution, training points appear only inside inner products:

Find $\alpha_1 \dots \alpha_N$ such that

$\mathbf{Q}(\boldsymbol{\alpha}) = \sum \alpha_i - \frac{1}{2} \sum \sum \alpha_i \alpha_j y_i y_j \mathbf{x}_i^T \mathbf{x}_j$ is maximized and

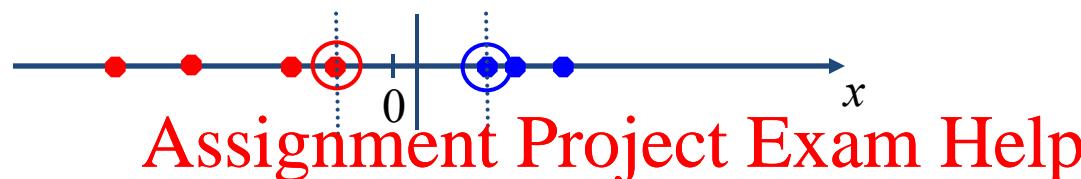
$$(1) \sum \alpha_i y_i = 0$$

$$(2) 0 \leq \alpha_i \leq C \text{ for all } \alpha_i$$

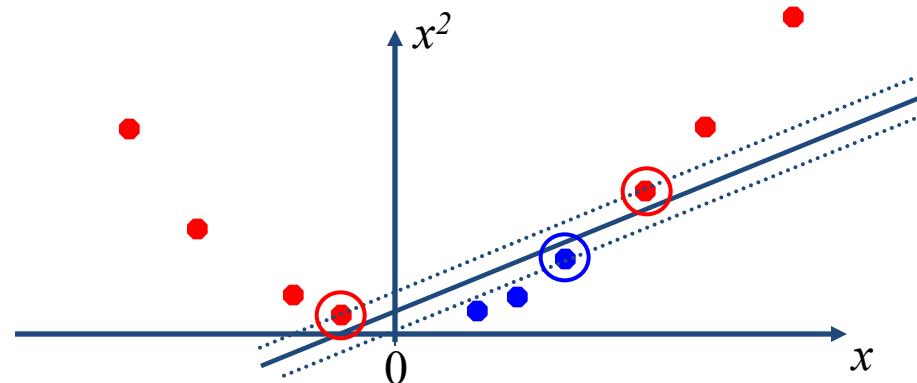
$$f(\mathbf{x}) = \sum_{sv \in SV} y_{sv} \cdot \alpha_{sv} \cdot \langle \mathbf{x}_{sv}, \mathbf{x} \rangle + b$$

Non-linear SVMs

- Datasets that are linearly separable (with some noise) work out great:

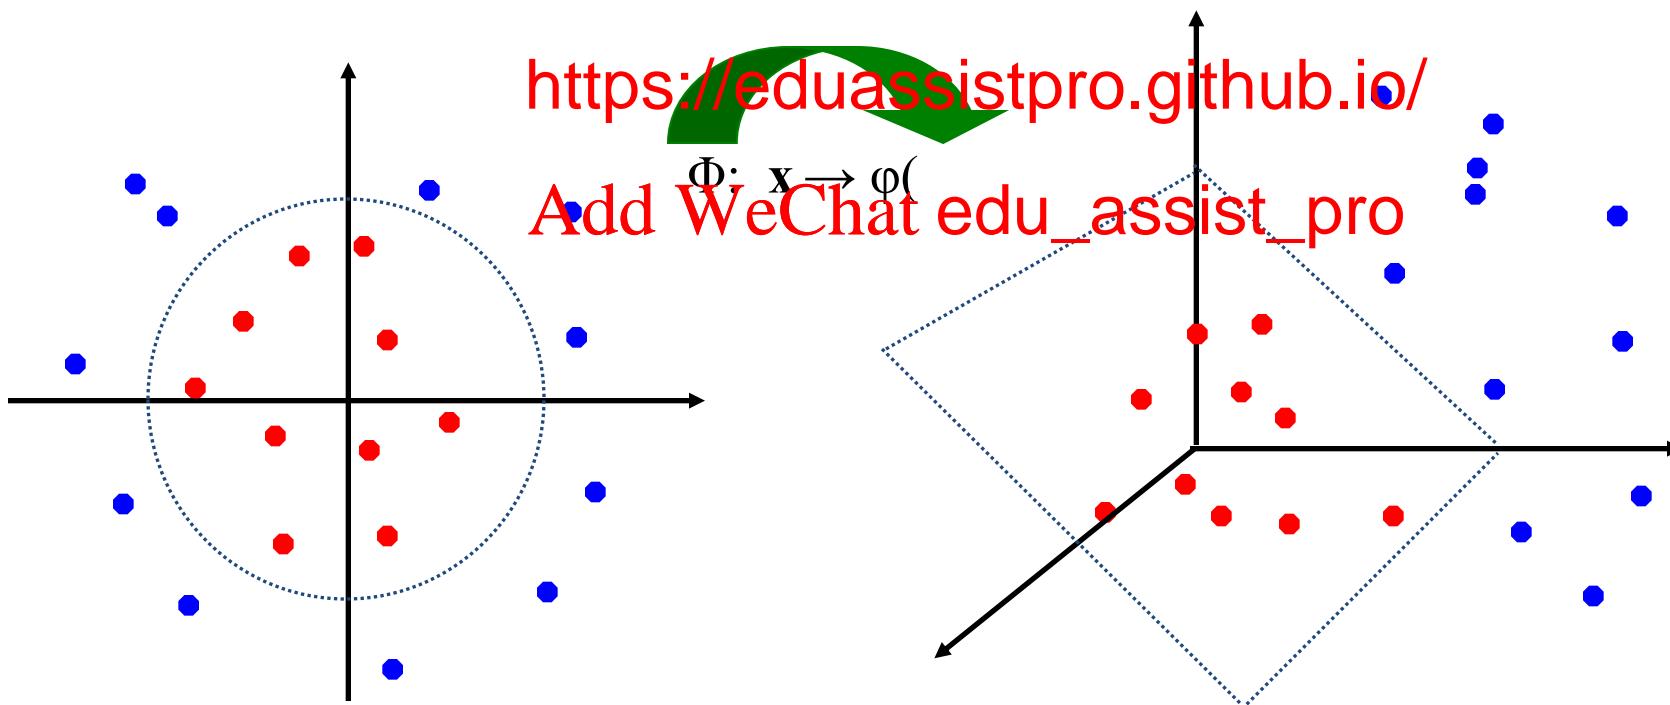


- But what are we going to do with non-linearly separable datasets?
- <https://eduassistpro.github.io/>
-
- A 1D plot on the x -axis. Red and blue dots are intermixed. A solid horizontal line represents the decision boundary. Several points are highlighted with circles: two red points on the left of the boundary and three blue points on the right of the boundary.
- Add WeChat [edu_assist_pro](https://eduassistpro.github.io/)
- How about ... mapping data to a higher-dimensional space:



Non-linear SVMs: Feature spaces

- General idea: the original feature space can always be mapped to some higher-dimensional feature space where the training set is separable:



The “Kernel Trick”

- The linear classifier relies on an inner product between vectors
 $K(\mathbf{x}_i, \mathbf{x}_j) = \mathbf{x}_i^T \mathbf{x}_j$
- If every datapoint is mapped into high-dimensional space via some transform, the inner product becomes:
<https://eduassistpro.github.io/>
- A *kernel function* is some function that maps to an inner product in **some** expanded feature space.
 - Usually, no need to construct the feature space explicitly.

What about $K(\mathbf{u}, \mathbf{v}) = (1 + \mathbf{u}^\top \mathbf{v})^3$?

Example

$$\begin{aligned} K(\mathbf{u}, \mathbf{v}) &= (1 + \mathbf{u}^\top \mathbf{v})^2 \\ &= 1 + 2\mathbf{u}^\top \mathbf{v} + (\mathbf{u}^\top \mathbf{v})^2 \end{aligned}$$

$$= 1 + 2\mathbf{u}^\top \mathbf{v} + \sum_{i=1}^d \mathbf{u}_i \mathbf{v}_i$$

$$\begin{aligned} &\text{Assignment Project Exam Help} \\ &= 1 + 2\mathbf{u}^\top \mathbf{v} + \sum_{i=1}^d \mathbf{u}_i^2 \mathbf{v}_i^2 + \sum_{i \neq j} \mathbf{u}_i \mathbf{v}_j \\ &= \phi(\mathbf{u})^\top \phi(\mathbf{v}) \end{aligned}$$

<https://eduassistpro.github.io/>
 Add WeChat edu_assist_pro

$$\phi(\mathbf{u}) = [1 \quad \sqrt{2}\mathbf{u}_1 \quad \dots \quad \sqrt{2}\mathbf{u}_d \quad \mathbf{u}_1^2 \quad \dots \quad \mathbf{u}_d^2 \quad \mathbf{u}_1 \mathbf{u}_2 \quad \dots \quad \mathbf{u}_{d-1} \mathbf{u}_d]^\top$$

$$\phi(\mathbf{v}) = [1 \quad \sqrt{2}\mathbf{v}_1 \quad \dots \quad \sqrt{2}\mathbf{v}_d \quad \mathbf{v}_1^2 \quad \dots \quad \mathbf{v}_d^2 \quad \mathbf{v}_1 \mathbf{v}_2 \quad \dots \quad \mathbf{v}_{d-1} \mathbf{v}_d]^\top$$

Linear

Non-linear

Non-linear + feature combination

Kernels

- Why use kernels?
 - Make non-separable problem separable.
 - Map data into better representational space
 - Can be learned to capture some notion of “similarity”
- Common kern <https://eduassistpro.github.io/>
 - Linear
 - Polynomial $K(\mathbf{x}, \mathbf{z}) = (\mathbf{1} + \mathbf{x}^T \mathbf{z})^d$
Add WeChat edu_assist_pro
 - Gives feature combinations
 - Radial basis function (infinite dimensional space)
$$K(\mathbf{x}_i, \mathbf{x}_j; \sigma) = \exp\left(-\frac{\|\mathbf{x}_i - \mathbf{x}_j\|^2}{2\sigma^2}\right)$$
- Text classification:
 - Usually use linear or polynomial kernels

Classification with SVMs **with Kernels**

- Given a new point \mathbf{x} , we can compute its score

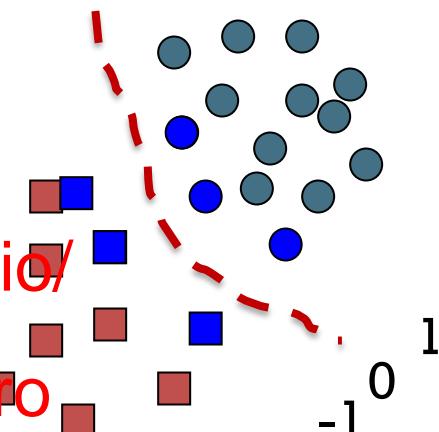
i.e.,

$$\sum_{sv \in SV} \alpha_{sv} K(\mathbf{x}_{sv}, \mathbf{x}) + b$$

- Decide c_l <https://eduassistpro.github.io/>

- E.g., in scikit-learn

Add WeChat edu_assist_pro



- linear: $\langle x, x' \rangle$.
- polynomial: $(\gamma \langle x, x' \rangle + r)^d$. d is specified by keyword `degree`, r by `coef0`.
- rbf: $\exp(-\gamma \|x - x'\|^2)$. γ is specified by keyword `gamma`, must be greater than 0.
- sigmoid ($\tanh(\gamma \langle x, x' \rangle + r)$), where r is specified by `coef0`.

Pros and Cons of the SVM Classifier

- Pros
 - High accuracy
 - Fast classification speed
 - Works with c samples <https://eduassistpro.github.io/> Kernel)
 - Can adapt to
 - Any $K(u, v)$ can be used if symmetric, and positive semi-definite.
 - Or explicitly engineer the feature space.
- Cons
 - Training does not scale well with number of training samples ($O(d*n^2)$ or $O(d*n^3)$)
 - Hyper-parameters need tuning

Resources for today's lecture

- Christopher J. C. Burges. 1998. A Tutorial on Support Vector Machines for Pattern Recognition
- S. T. Dumais. 1998. Using SVMs for text categorization, IEEE Intelligent Systems, 13(4)
- S. T. Dumais, J. Platt, D. Heckerman and M. Sahami. 1998. Inductive learning algorithms and representations for text categorization. *CIKM '98*, pp. 148-155.
- Yiming Yang, Xin Liu. 1999. A re-examination of text categorization methods. 22nd Annual International SIGIR
- Tong Zhang, Frank J. Ole <https://eduassistpro.github.io/> Regularized Linear Classification Methods. I
- Trevor Hastie, Robert Tibshirani and Jerome Friedman *Statistical Learning: Data Mining, Inference and Prediction*. Springer-Verlag
- T. Joachims, *Learning to Classify Text using Support Vector Machines*. Kluwer, 2002.
- Fan Li, Yiming Yang. 2003. A Loss Function Analysis for Classification Methods in Text Categorization. ICML 2003: 472-479.
- Tie-Yan Liu, Yiming Yang, Hao Wan, et al. 2005. Support Vector Machines Classification with Very Large Scale Taxonomy, SIGKDD Explorations, 7(1): 36-43.
- ‘Classic’ Reuters-21578 data set: <http://www.daviddlewis.com/resources/testcollections/reuters21578/>